

report

November 9, 2024

1 Report

Let's set the random seeds for reproducibility.

```
[1]: import numpy as np
import random
import torch
import lightning as L

RANDOM_SEED = 42
random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)
torch.manual_seed(RANDOM_SEED)
torch.cuda.manual_seed_all(RANDOM_SEED)
L.seed_everything(RANDOM_SEED)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
```

Seed set to 42

1.1 Data Preprocessing

We will start by loading the data for the Apple stock and inspecting their first few rows.

```
[2]: import yfinance as yf

# Download Apple stock data
stock_prices = yf.download('AAPL',
                           start='2015-01-01',
                           end='2024-01-31')

# Display the first few rows
print("First few rows of Apple stock data:")
stock_prices.head()
```

```
[*****100%*****] 1 of 1 completed
```

First few rows of Apple stock data:

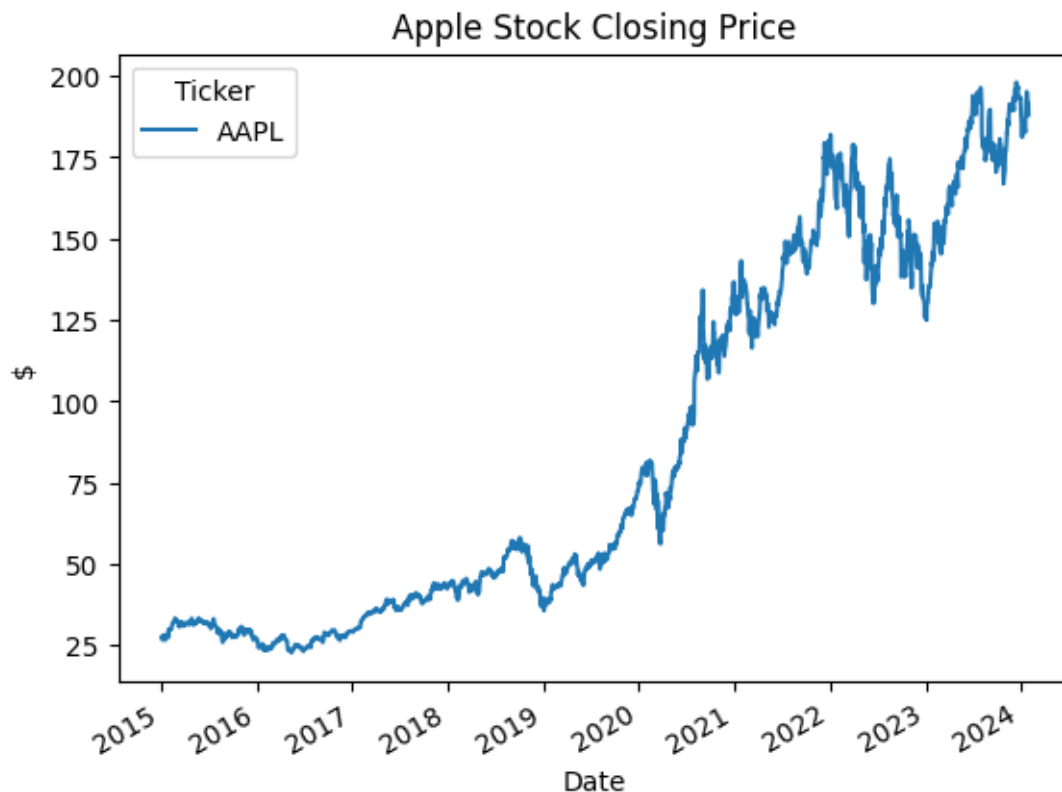
```
[2]: Price                Adj Close      Close      High      Low \
      Ticker                AAPL        AAPL        AAPL        AAPL
      Date
      2015-01-02 00:00:00+00:00  24.373955  27.332500  27.860001  26.837500
      2015-01-05 00:00:00+00:00  23.687304  26.562500  27.162500  26.352501
      2015-01-06 00:00:00+00:00  23.689529  26.565001  26.857500  26.157499
      2015-01-07 00:00:00+00:00  24.021709  26.937500  27.049999  26.674999
      2015-01-08 00:00:00+00:00  24.944679  27.972500  28.037500  27.174999

      Price                Open      Volume
      Ticker                AAPL        AAPL
      Date
      2015-01-02 00:00:00+00:00  27.847500  212818400
      2015-01-05 00:00:00+00:00  27.072500  257142000
      2015-01-06 00:00:00+00:00  26.635000  263188400
      2015-01-07 00:00:00+00:00  26.799999  160423600
      2015-01-08 00:00:00+00:00  27.307501  237458000
```

Let's examine the path of the stock's closing price over time.

```
[3]: import matplotlib.pyplot as plt

stock_prices['Close'].plot()
plt.title('Apple Stock Closing Price')
plt.xlabel('Date')
plt.ylabel('$')
plt.show()
```



The stock price has a clear upward (multiplicative) trend over time as expected from such a tech company. There is a dip in early 2019 and 2020 probably due to the COVID-19 pandemic, which is followed by a significant growth rate. The volatility seems to be increasing as the price is increasing.

Let's check for any missing values in the dataset.

```
[4]: # Check for missing values in the dataset
stock_prices.isnull().sum()
```

```
[4]: Price      Ticker
Adj Close  AAPL      0
Close      AAPL      0
High       AAPL      0
Low        AAPL      0
Open       AAPL      0
Volume     AAPL      0
dtype: int64
```

There are none.

We proceed with the calculation of the daily returns. We expect to have only one missing value in the daily returns due to the nature of the calculation.

We then define the extreme events, which will act as our target variable.

```
[5]: # Calculate the daily returns
stock_prices['Daily_Returns'] = stock_prices['Adj Close'].pct_change() * 100

# Remove the first row as it is a NaN value
stock_prices = stock_prices.dropna()
stock_prices.head()
```

```
[5]: Price                Adj Close      Close      High      Low \
      Ticker                AAPL      AAPL      AAPL      AAPL
      Date
2015-01-05 00:00:00+00:00  23.687304  26.562500  27.162500  26.352501
2015-01-06 00:00:00+00:00  23.689529  26.565001  26.857500  26.157499
2015-01-07 00:00:00+00:00  24.021709  26.937500  27.049999  26.674999
2015-01-08 00:00:00+00:00  24.944679  27.972500  28.037500  27.174999
2015-01-09 00:00:00+00:00  24.971437  28.002501  28.312500  27.552500
```

```
Price                Open      Volume Daily_Returns
Ticker                AAPL      AAPL
Date
2015-01-05 00:00:00+00:00  27.072500  257142000      -2.817151
2015-01-06 00:00:00+00:00  26.635000  263188400       0.009397
2015-01-07 00:00:00+00:00  26.799999  160423600       1.402223
2015-01-08 00:00:00+00:00  27.307501  237458000       3.842232
2015-01-09 00:00:00+00:00  28.167500  214798000       0.107270
```

```
[6]: # Create the Extreme_Event column
stock_prices['Extreme_Event'] = (abs(stock_prices['Daily_Returns']) > 2).
    .astype(int)
stock_prices.head()
```

```
[6]: Price                Adj Close      Close      High      Low \
      Ticker                AAPL      AAPL      AAPL      AAPL
      Date
2015-01-05 00:00:00+00:00  23.687304  26.562500  27.162500  26.352501
2015-01-06 00:00:00+00:00  23.689529  26.565001  26.857500  26.157499
2015-01-07 00:00:00+00:00  24.021709  26.937500  27.049999  26.674999
2015-01-08 00:00:00+00:00  24.944679  27.972500  28.037500  27.174999
2015-01-09 00:00:00+00:00  24.971437  28.002501  28.312500  27.552500
```

```
Price                Open      Volume Daily_Returns Extreme_Event
Ticker                AAPL      AAPL
Date
2015-01-05 00:00:00+00:00  27.072500  257142000      -2.817151           1
2015-01-06 00:00:00+00:00  26.635000  263188400       0.009397           0
2015-01-07 00:00:00+00:00  26.799999  160423600       1.402223           0
2015-01-08 00:00:00+00:00  27.307501  237458000       3.842232           1
2015-01-09 00:00:00+00:00  28.167500  214798000       0.107270           0
```

We expect that the dataset will be imbalanced. Let's verify that.

```
[7]: # Count the number of extreme events
extreme_events = stock_prices[stock_prices['Extreme_Event'] == 1]
print(f"Number of extreme events: {len(extreme_events)}")

# Print the percentage of extreme events and no extreme events in the dataset
print(f"Percentage of extreme events: {len(extreme_events) / len(stock_prices) * 100:.2f}%")
print(f"Percentage of no extreme events: {(len(stock_prices) - len(extreme_events)) / len(stock_prices) * 100:.2f}%")
```

```
Number of extreme events: 460
Percentage of extreme events: 20.15%
Percentage of no extreme events: 79.85%
```

The dataset is moderately imbalanced with approximately 20% of the data being extreme events. This is a challenge that we will have to address in the modelling phase.

Our choices for doing so are limited, since we are dealing with a time-series dataset and we need to preserve the temporal nature of the data, which excludes the traditional use of resampling techniques, such as upsampling the minority class or undersampling the majority class.

Nevertheless, we will mitigate this issue by using weighted loss functions for the neural networks and balanced class weights for the random forest.

Next, we shift the target variable `Extreme_Event` by one day so that the model is trained to predict if an extreme event occurs **tomorrow** based on today's data.

```
[8]: # Shift the Extreme_Event column by one day
stock_prices['Extreme_Event'] = stock_prices['Extreme_Event'].shift(-1)

# Remove the last row as it is a NaN value
stock_prices = stock_prices.dropna()
stock_prices['Extreme_Event'] = stock_prices['Extreme_Event'].astype(int)
stock_prices.head()
```

```
[8]: Price           Adj Close      Close      High      Low \
      Ticker           AAPL      AAPL      AAPL      AAPL
      Date
2015-01-05 00:00:00+00:00  23.687304  26.562500  27.162500  26.352501
2015-01-06 00:00:00+00:00  23.689529  26.565001  26.857500  26.157499
2015-01-07 00:00:00+00:00  24.021709  26.937500  27.049999  26.674999
2015-01-08 00:00:00+00:00  24.944679  27.972500  28.037500  27.174999
2015-01-09 00:00:00+00:00  24.971437  28.002501  28.312500  27.552500

      Price           Open      Volume  Daily_Returns  Extreme_Event
      Ticker           AAPL      AAPL
      Date
```

2015-01-05 00:00:00+00:00	27.072500	257142000	-2.817151	0
2015-01-06 00:00:00+00:00	26.635000	263188400	0.009397	0
2015-01-07 00:00:00+00:00	26.799999	160423600	1.402223	1
2015-01-08 00:00:00+00:00	27.307501	237458000	3.842232	0
2015-01-09 00:00:00+00:00	28.167500	214798000	0.107270	1

We continue by splitting the data into training, validation and test sets.

```
[9]: # Extract features and labels
features = stock_prices[['Open', 'High', 'Low', 'Close', 'Volume',
↪ 'Daily>Returns']]
labels = stock_prices['Extreme_Event'].astype(int) # Target variable

# Define train/val/test split ratios
train_ratio = 0.7
val_ratio = 0.85

# Split into train, validation and test sets
train_size = int(train_ratio * len(stock_prices))
val_size = int(val_ratio * len(stock_prices))

X_train = features.iloc[:train_size]
y_train = labels.iloc[:train_size]

X_val = features.iloc[train_size:val_size]
y_val = labels.iloc[train_size:val_size]

X_test = features.iloc[val_size:]
y_test = labels.iloc[val_size:]

print(f"Training features shape: {X_train.shape}, labels shape: {y_train.
↪ shape}")
print(f"Validation features shape: {X_val.shape}, labels shape: {y_val.shape}")
print(f"Test features shape: {X_test.shape}, labels shape: {y_test.shape}")
```

Training features shape: (1597, 6), labels shape: (1597,)

Validation features shape: (342, 6), labels shape: (342,)

Test features shape: (343, 6), labels shape: (343,)

Let's check the distribution of labels in the training, validation and test sets. We want to avoid overestimating the performance of the model on unseen data, due to total lack of positive examples.

```
[10]: import matplotlib.pyplot as plt
import numpy as np

# Calculate distributions
train_dist = y_train.value_counts() / len(y_train)
val_dist = y_val.value_counts() / len(y_val)
```

```

test_dist = y_test.value_counts() / len(y_test)

# Create bar plot
fig, ax = plt.subplots(figsize=(10, 6))

x = np.arange(3)
width = 0.35

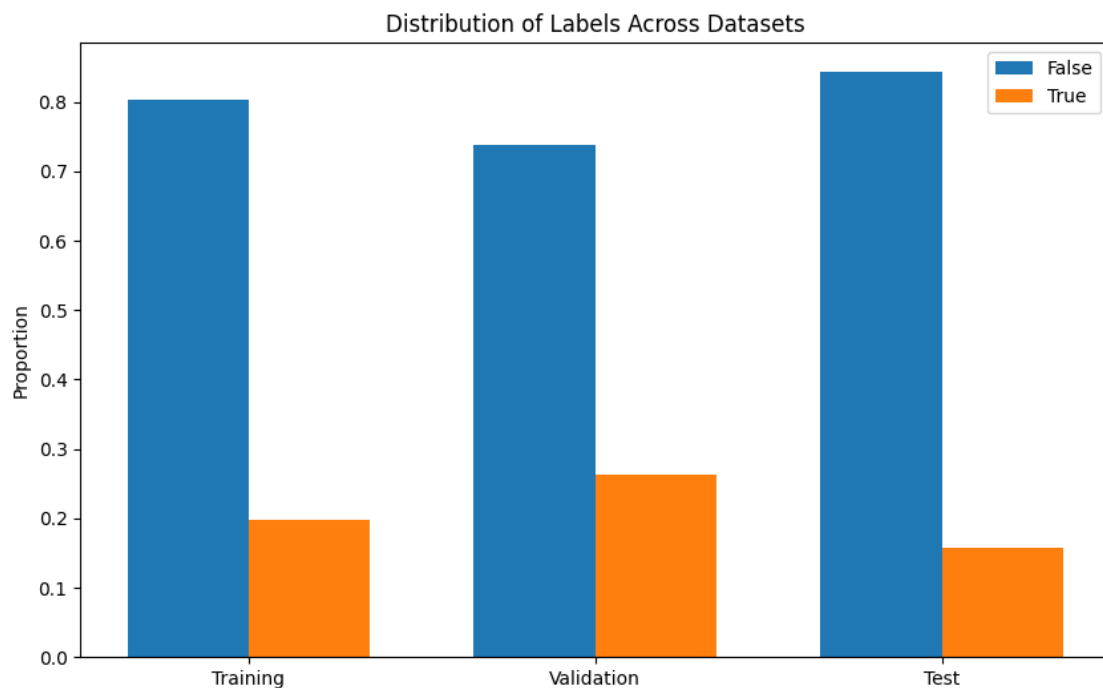
# Group bars by False/True
false_vals = [train_dist[0], val_dist[0], test_dist[0]]
true_vals = [train_dist[1], val_dist[1], test_dist[1]]

ax.bar(x - width/2, false_vals, width, label='False')
ax.bar(x + width/2, true_vals, width, label='True')

ax.set_xticks(x)
ax.set_xticklabels(['Training', 'Validation', 'Test'])
ax.set_ylabel('Proportion')
ax.set_title('Distribution of Labels Across Datasets')
ax.legend()

plt.show()

```



The distribution of labels is similar across the datasets, so we can proceed with fitting the models.

As a final step of preprocessing, we will transform the data into a more appropriate format for supervised learning. Specifically, for each day, we will create a sequence of historical features by going back 10 days. Hence, since the dataset has 6 features, we will end up with 60 features for each day. The corresponding label will be 1 if an extreme event occurs tomorrow and 0 otherwise.

```
[11]: from data_preprocessing import StockDataPreprocessor

X_train, y_train = StockDataPreprocessor.time_series_to_supervised(
    X_train, y_train, lookback=10
)
X_val, y_val = StockDataPreprocessor.time_series_to_supervised(X_val, y_val,
    ↪lookback=10)
X_test, y_test = StockDataPreprocessor.time_series_to_supervised(X_test,
    ↪y_test, lookback=10)

print(f"Training features shape: {X_train.shape}, labels shape: {y_train.
    ↪shape}")
print(f"Validation features shape: {X_val.shape}, labels shape: {y_val.shape}")
print(f"Test features shape: {X_test.shape}, labels shape: {y_test.shape}")
```

Training features shape: (1587, 60), labels shape: (1587,)

Validation features shape: (332, 60), labels shape: (332,)

Test features shape: (333, 60), labels shape: (333,)

Note: All the above steps are implemented in the `src/data_preprocessing.py` file.

Let's examine the first few rows of the training set. For each feature, we will have 10 days of history, amounting to 60 features in total.

```
[12]: print(f'Shape of the training set: {X_train.shape}')
X_train.head()
```

Shape of the training set: (1587, 60)

```
[12]:
```

	Open(t-10)	High(t-10)	Low(t-10)	Close(t-10)	\
Date					
2015-01-20 00:00:00+00:00	27.072500	27.162500	26.352501	26.562500	
2015-01-21 00:00:00+00:00	26.635000	26.857500	26.157499	26.565001	
2015-01-22 00:00:00+00:00	26.799999	27.049999	26.674999	26.937500	
2015-01-23 00:00:00+00:00	27.307501	28.037500	27.174999	27.972500	
2015-01-26 00:00:00+00:00	28.167500	28.312500	27.552500	28.002501	

	Volume(t-10)	Daily_Returns(t-10)	Open(t-9)	\
Date				
2015-01-20 00:00:00+00:00	257142000.0	-2.817151	26.635000	
2015-01-21 00:00:00+00:00	263188400.0	0.009397	26.799999	
2015-01-22 00:00:00+00:00	160423600.0	1.402223	27.307501	
2015-01-23 00:00:00+00:00	237458000.0	3.842232	28.167500	
2015-01-26 00:00:00+00:00	214798000.0	0.107270	28.150000	

Date	High(t-9)	Low(t-9)	Close(t-9)	...	Low(t-2)	\
2015-01-20 00:00:00+00:00	26.857500	26.157499	26.565001	...	26.665001	
2015-01-21 00:00:00+00:00	27.049999	26.674999	26.937500	...	26.299999	
2015-01-22 00:00:00+00:00	28.037500	27.174999	27.972500	...	26.625000	
2015-01-23 00:00:00+00:00	28.312500	27.552500	28.002501	...	27.067499	
2015-01-26 00:00:00+00:00	28.157499	27.200001	27.312500	...	27.430000	

Date	Close(t-2)	Volume(t-2)	Daily_Returns(t-2)	\
2015-01-20 00:00:00+00:00	26.705000	240056000.0	-2.714021	
2015-01-21 00:00:00+00:00	26.497499	314053200.0	-0.777047	
2015-01-22 00:00:00+00:00	27.180000	199599600.0	2.575728	
2015-01-23 00:00:00+00:00	27.387501	194303600.0	0.763436	
2015-01-26 00:00:00+00:00	28.100000	215185600.0	2.601532	

Date	Open(t-1)	High(t-1)	Low(t-1)	Close(t-1)	\
2015-01-20 00:00:00+00:00	26.757500	26.895000	26.299999	26.497499	
2015-01-21 00:00:00+00:00	26.959999	27.242500	26.625000	27.180000	
2015-01-22 00:00:00+00:00	27.237499	27.764999	27.067499	27.387501	
2015-01-23 00:00:00+00:00	27.565001	28.117500	27.430000	28.100000	
2015-01-26 00:00:00+00:00	28.075001	28.437500	27.882500	28.245001	

Date	Volume(t-1)	Daily_Returns(t-1)
2015-01-20 00:00:00+00:00	314053200.0	-0.777047
2015-01-21 00:00:00+00:00	199599600.0	2.575728
2015-01-22 00:00:00+00:00	194303600.0	0.763436
2015-01-23 00:00:00+00:00	215185600.0	2.601532
2015-01-26 00:00:00+00:00	185859200.0	0.516030

[5 rows x 60 columns]

2 Random Forest Model

We now proceed with fitting the random forest model. Since we are using a tree-based model, we will not normalize the data.

In order to tune the hyperparameters of the random forest model, we will use the **hyperopt** library. It uses a Bayesian optimization approach to find the best hyperparameters, called **Tree of Parzen Estimators** (TPE). This method models the likelihood of the objective function, given the data and a choice of hyperparameters, using kernel density estimation.

The choice of metric for the optimization is nuanced when dealing with imbalanced datasets and one should consider what is the cost of the false negatives and false positives. In our case, a false positive would mean that we inaccurately predict an extreme event for the next day, which could

lead to a loss of capital and trust in the model. On the other hand, a false negative would mean that we inaccurately predict no extreme event for the next day, which could lead to a loss of potential profits.

Precision is preferred when false positives are more costly and recall is preferred when false negatives are more costly. Assuming that we are not highly risk averse, i.e. we are willing to lose some money in order to make a profit, we will use the metric F_β with $\beta = 2$ as our optimization metric, which weighs recall more heavily than precision.

F_2 is defined as:

$$F_2 = (1 + 2^2) \cdot \frac{\text{Precision} \cdot \text{Recall}}{2^2 \cdot \text{Precision} + \text{Recall}},$$

where $\text{Precision} = \frac{TP}{TP+FP}$ and $\text{Recall} = \frac{TP}{TP+FN}$.

```
[13]: from hyperopt import fmin, tpe, hp, STATUS_OK, Trials
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import fbeta_score

      def objective(params):
          """
          Objective function for hyperopt optimization.
          """
          model_params = {
              'n_estimators': int(params['n_estimators']),
              'max_depth': int(params['max_depth']),
              'min_samples_split': int(params['min_samples_split']),
              'min_samples_leaf': int(params['min_samples_leaf']),
              'max_leaf_nodes': int(params['max_leaf_nodes']),
              'max_features': params['max_features'],
              'min_impurity_decrease': params['min_impurity_decrease'],
              'criterion': 'entropy',
              'class_weight': 'balanced',
              'bootstrap': True,
              'random_state': 42,
              'n_jobs': -1
          }

          # Create and train model
          model = RandomForestClassifier(**model_params)
          model.fit(X_train.values, y_train.values.astype(int))

          # Evaluate on validation set
          y_pred = model.predict(X_val.values)
          f2 = fbeta_score(y_val.values.astype(int), y_pred, beta=2)

          return {'loss': -f2, 'status': STATUS_OK}

      # Define search space
```

```

space = {
    'n_estimators': hp.quniform('n_estimators', 700, 800, 1),
    'max_depth': hp.quniform('max_depth', 70, 100, 1),
    'min_samples_split': hp.quniform('min_samples_split', 5, 15, 1),
    'min_samples_leaf': hp.quniform('min_samples_leaf', 1, 3, 1),
    'max_features': hp.uniform('max_features', 0.8, 1.0),
    'max_leaf_nodes': hp.quniform('max_leaf_nodes', 500, 600, 1),
    'min_impurity_decrease': hp.uniform('min_impurity_decrease', 0.04, 0.05),
}

# Run optimization
trials = Trials()
best = fmin(fn=objective,
            space=space,
            algo=tpe.suggest,
            max_evals=30,
            trials=trials)

print("\nBest trial:")
print(f" Value (F2): {-min(trials.losses()):.4f}")
print("\nBest parameters:")
for key, value in best.items():
    print(f"    {key}: {value}")

```

```

100%|      | 30/30 [00:34<00:00, 1.16s/trial, best loss:
-0.6432748538011696]

```

```

Best trial:
Value (F2): 0.6433

```

```

Best parameters:
max_depth: 96.0
max_features: 0.8098582865554091
max_leaf_nodes: 599.0
min_impurity_decrease: 0.04380604831017937
min_samples_leaf: 2.0
min_samples_split: 12.0
n_estimators: 722.0

```

We proceed with fitting the final model with the best parameters and evaluating on the test set.

```

[14]: # Train final model with best parameters
best_params = best.copy()
best_params['n_estimators'] = int(best_params['n_estimators'])
best_params['max_depth'] = int(best_params['max_depth'])
best_params['min_samples_split'] = int(best_params['min_samples_split'])
best_params['min_samples_leaf'] = int(best_params['min_samples_leaf'])
best_params['max_leaf_nodes'] = int(best_params['max_leaf_nodes'])

```

```

best_params.update({
    'criterion': 'entropy',
    'class_weight': 'balanced',
    'bootstrap': True,
    'random_state': 42,
    'n_jobs': -1
})

final_model = RandomForestClassifier(**best_params)
final_model.fit(X_train.values, y_train.values.astype(int))

```

```

[14]: RandomForestClassifier(class_weight='balanced', criterion='entropy',
                             max_depth=96,
                             max_features=np.float64(0.8098582865554091),
                             max_leaf_nodes=599,
                             min_impurity_decrease=np.float64(0.04380604831017937),
                             min_samples_leaf=2, min_samples_split=12,
                             n_estimators=722, n_jobs=-1, random_state=42)

```

```

[15]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import (accuracy_score, roc_auc_score,
    ↪balanced_accuracy_score,
                                precision_score, recall_score, confusion_matrix,
    ↪f1_score)

# Make predictions
y_pred = final_model.predict(X_test.values)
y_true = y_test.values.astype(int)

# Calculate metrics
metrics = {
    'F2 Score': fbeta_score(y_true, y_pred, beta=2),
    'F1 Score': fbeta_score(y_true, y_pred, beta=1),
    'Precision': precision_score(y_true, y_pred),
    'Recall': recall_score(y_true, y_pred),
    'Accuracy': accuracy_score(y_true, y_pred),
    'AUC': roc_auc_score(y_true, y_pred),
    'Balanced Accuracy': balanced_accuracy_score(y_true, y_pred)
}

# Print metrics in a organized way
print("Test Set Metrics:")
print("-" * 40)
for metric, value in metrics.items():
    print(f"{metric:20s}: {value:.4f}")

```

```

print("-" * 40)

# Plot confusion matrix using seaborn
plt.figure(figsize=(6, 4))
cm = confusion_matrix(y_true, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Event', 'Extreme Event'],
            yticklabels=['No Event', 'Extreme Event'])
plt.title('Confusion Matrix')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.tight_layout()
plt.show()

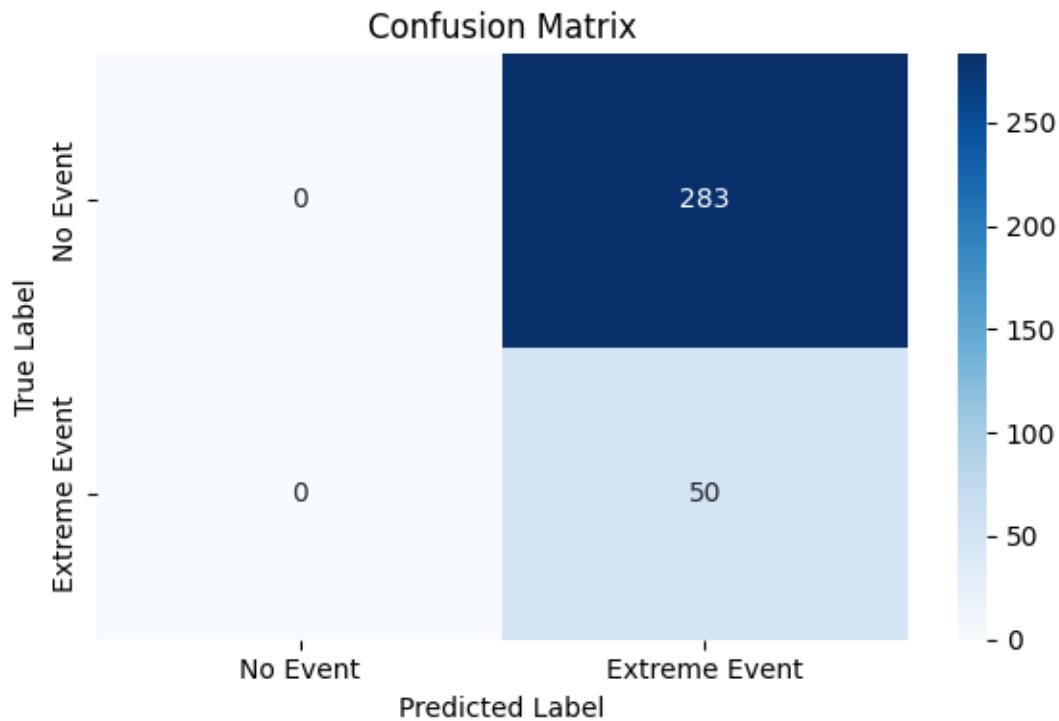
```

Test Set Metrics:

```

-----
F2 Score      : 0.4690
F1 Score      : 0.2611
Precision     : 0.1502
Recall        : 1.0000
Accuracy      : 0.1502
AUC           : 0.5000
Balanced Accuracy : 0.5000
-----

```



We see that the model is predicting all the test samples as extreme events, because of the metric we used for optimization. It has not learned how to truly distinguish the two classes, hence we are severely underfitting the training data.

3 Temporal CNN Model

Let's prepare the data for the temporal CNN model. Firstly, we will reload the data and then standardize the features, since we will be using a neural network. The daily returns do not need to be standardized due to their scale. For this purpose, we will use the `StockDataPreprocessor` class defined in the `src/data_preprocessing.py` file.

```
[2]: from data_preprocessing import StockDataPreprocessor

stock_prices = StockDataPreprocessor(
    ticker="AAPL", start_date='2015-01-01', end_date='2024-01-31'
).download_and_prepare_stock_data()
# Standardize the features
stock_prices = StockDataPreprocessor.standardize_data(stock_prices, ['Open', 'Adj_Close', 'High', 'Low', 'Close', 'Volume'])
stock_prices.head()
```

[*****100%*****] 1 of 1 completed

```
[2]: Price           Adj Close      Close      High      Low      Open \
      Ticker           AAPL        AAPL        AAPL        AAPL        AAPL
      Date
2015-01-05 00:00:00+00:00  23.661276 -1.031309 -1.025440 -1.030020 -1.022063
2015-01-06 00:00:00+00:00  23.663504 -1.031264 -1.030777 -1.033506 -1.029804
2015-01-07 00:00:00+00:00  23.995312 -1.024681 -1.027408 -1.024256 -1.026884
2015-01-08 00:00:00+00:00  24.917271 -1.006389 -1.010130 -1.015318 -1.017906
2015-01-09 00:00:00+00:00  24.943998 -1.005859 -1.005318 -1.008570 -1.002691
```

```
Price           Volume Daily_Returns Extreme_Event
Ticker           AAPL
Date
2015-01-05 00:00:00+00:00  1.965753      -2.817141          0.0
2015-01-06 00:00:00+00:00  2.054466       0.009415          0.0
2015-01-07 00:00:00+00:00  0.546695       1.402193          1.0
2015-01-08 00:00:00+00:00  1.676948       3.842246          0.0
2015-01-09 00:00:00+00:00  1.344479       0.107266          1.0
```

We continue by splitting the data into training, validation and test sets.

```
[3]: # Split data into train, validation and test sets
X_train, y_train, X_val, y_val, X_test, y_test = StockDataPreprocessor.
    ↪split_data()
```

```

stock_prices,
["Open", "High", "Low", "Close", "Volume", "Daily_Returns"],
"Extreme_Event",
train_ratio=0.7,
val_ratio=0.85,
)

```

Finally, to predict whether an extreme event will occur tomorrow, given the past 10 days, we ought to transform the data into sequences with shape `[n_samples, n_features, lookback]`, where `lookback = 10`.

```

[4]: X_train, y_train = StockDataPreprocessor.create_sequences(X_train, y_train,
    ↪lookback=10)
X_val, y_val = StockDataPreprocessor.create_sequences(X_val, y_val, lookback=10)
X_test_tcnn, y_test_tcnn = StockDataPreprocessor.create_sequences(X_test,
    ↪y_test, lookback=10)

print(f"Training features shape: {X_train.shape}, labels shape: {y_train.
    ↪shape}")
print(f"Validation features shape: {X_val.shape}, labels shape: {y_val.shape}")
print(f"Test features shape: {X_test.shape}, labels shape: {y_test.shape}")

```

```

Training features shape: (1587, 6, 10), labels shape: (1587,)
Validation features shape: (332, 6, 10), labels shape: (332,)
Test features shape: (343, 6), labels shape: (343,)

```

For the definition of the temporal CNN model and the management of the training process, we will use the `lightning` library. All the relevant code is implemented in the `src/temporal_cnn.py` file. We import the necessary functions and classes, we train the model and plot the training and validation losses.

The AdamW optimizer is chosen, which is a combination of Adam and Tikhonov regularization, to prevent overfitting. In addition, the learning rate is automatically adjusted using the ReduceLROnPlateau scheduler. Finally, the training includes early stopping and model checkpointing callbakcs.

```

[5]: from temporal_cnn import TCNNClassifier

tcnn_model = TCNNClassifier(n_features=X_train.shape[1],
    lookback=X_train.shape[2],
    hidden_dim=64,
    conv_channels=96,
    kernel_size=3,
    dropout_prob=0.3,
    learning_rate=1e-4,
    scheduler_patience=5,
    scheduler_factor=0.5,
    min_lr=1e-6)

```

```
tcnn_model.train(
    X_train, y_train, X_val, y_val, batch_size=64, max_epochs=100,
    patience=10
)
tcnn_model.plot_training_history()
```

Seed set to 42

GPU available: True (mps), used: True

TPU available: False, using: 0 TPU cores

HPU available: False, using: 0 HPUs

/Users/nikolaosmourdoukoutas/ai2c_assignment/.venv/lib/python3.10/site-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

warnings.warn(

	Name	Type	Params	Mode
0	cnn	Sequential	131 K	train

131 K	Trainable params
0	Non-trainable params
131 K	Total params
0.525	Total estimated model params size (MB)
12	Modules in train mode
0	Modules in eval mode

Sanity Checking: | | 0/? [00:00<?, ?it/s]

/Users/nikolaosmourdoukoutas/ai2c_assignment/.venv/lib/python3.10/site-packages/lightning/pytorch/trainer/connectors/data_connector.py:424: The 'val_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=9` in the `DataLoader` to improve performance.

/Users/nikolaosmourdoukoutas/ai2c_assignment/.venv/lib/python3.10/site-packages/lightning/pytorch/trainer/connectors/data_connector.py:424: The 'train_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=9` in the `DataLoader` to improve performance.

Training: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]


Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]

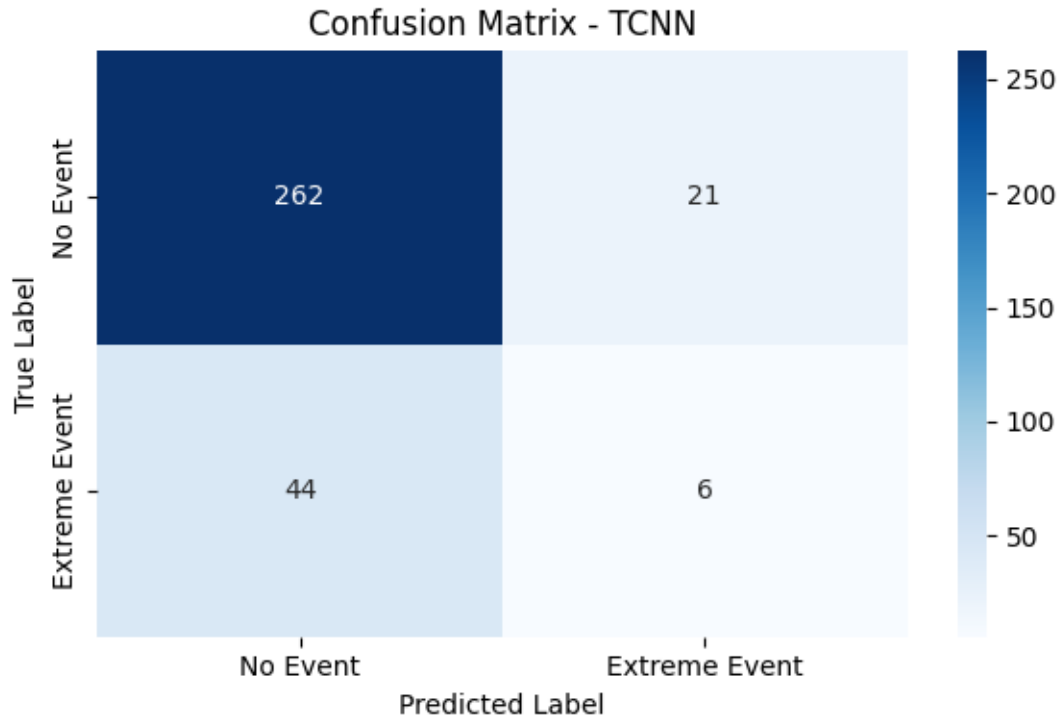


We see that the model is far from converged, since both losses are fluctuating and not decreasing. Let's load the best model and evaluate it on the test set.

```
[26]: from temporal_cnn import TCNN
      from model_evaluation import ModelEvaluator
      tcnn_model = TCNN(
          n_features=X_test_tcnn.shape[1],
          lookback=10,
          hidden_dim=128,
          conv_channels=64,
          kernel_size=3,
          dropout_prob=0.3,
          learning_rate=1e-4,
      )
      tcnn_model.load_state_dict(torch.load("../data/models/best_tcnn.pth"))
      evaluator = ModelEvaluator(model=tcnn_model, model_type="TCNN")
      metrics = evaluator.evaluate(X_test_tcnn, y_test_tcnn)
      for metric, value in metrics.items():
          print(f"{metric:20s}: {value:.4f}")
      print("-" * 40)

      evaluator.plot_confusion_matrix(X_test_tcnn, y_test_tcnn)
```

```
F2 Score           : 0.1322
F1 Score           : 0.1558
Precision          : 0.2222
Recall            : 0.1200
Accuracy          : 0.8048
AUC               : 0.5229
Balanced Accuracy  : 0.5229
```



The TCNN seems to be better than the random forest model, since it is able to accurately predict 6 extreme events and most of the non-extreme events. Its AUC and balanced accuracy scores are 3% higher than those of the random forest model. Finally, it is predicting accurately the majority of the non-extreme events. Hence, it has slightly learned to distinguish the two classes.

Let's plot the actual and predicted extreme events over time. We will mark the actual extreme events with blue stems and the predicted extreme events with red markers.

```
[27]: import matplotlib.pyplot as plt
import pandas as pd

y_test = y_test.iloc[10:]
y_pred = evaluator.predict(X_test_tcnn)
y_pred = pd.Series(y_pred, index=y_test.index)

plt.figure(figsize=(15, 6))

# Plot actual events as stems
plt.stem(y_test.index[y_test == 1], y_test[y_test == 1],
        label='Actual Events', linefmt='b-', markerfmt='bo', basefmt=' ')

# Plot predicted events as red markers
plt.plot(y_test.index[y_pred == 1], y_pred[y_pred == 1],
        'r^', label='Predicted Events', markersize=8)
```

```

plt.title('Extreme Events: Actual vs Predicted (TCNN)', pad=20)
plt.xlabel('Date')
plt.ylabel('Event Occurrence')

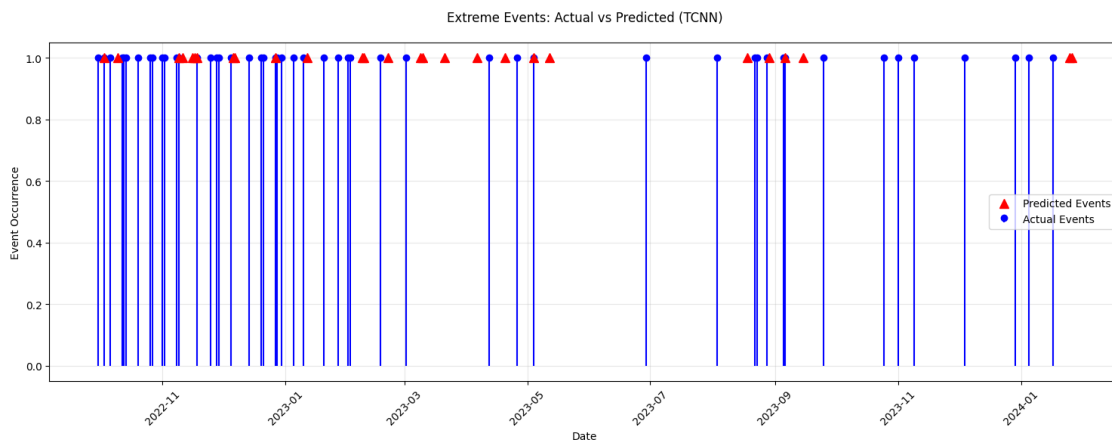
# Customize the plot
plt.grid(True, alpha=0.3)
plt.legend()

# Rotate x-axis labels for better readability
plt.xticks(rotation=45)

# Adjust layout to prevent label cutoff
plt.tight_layout()

# Save the plot
plt.savefig('../data/figures/tcnn_predictions.png', dpi=300,
            bbox_inches='tight')

```



The TCNN model demonstrates somewhat moderate prediction behavior, successfully identifying a few extreme events while maintaining relatively few false positives. The model's predictions show good alignment with actual events in the earlier period through March 2023, though its performance appears to decline after May 2023 where it misses most of the actual events. Still, it manages to predict a couple of extreme events after May 2023, where the true positive examples become sparser overall. It is possible that the fit of the LSTM would significantly improve, if we had a larger sample size or more balanced classes, providing a very good model for the task, since it would be to balance the precision and recall.

There are still several instances where actual events occur without corresponding predictions, suggesting room for improvement in the model's recall.

4 Improvements

4.1 Adding more features

Since we are interested in predicting extreme events, we will include five additional features that could provide useful information for the task at hand. Specifically, we will include:

- 10-day Rolling Volatility: Periods of high volatility often cluster together, and extreme events are more likely during volatile periods.
- Volume Relative to 10-day Moving Average: Unusual trading volume often precedes or accompanies extreme price movements.
- VIX Index: Market-wide volatility often correlates with individual stock extreme movements.
- Bollinger Band Width: Measures volatility expansion/contraction, particularly useful for identifying potential breakout periods.
- Average True Range (ATR): Captures true volatility including gaps, which is especially relevant for Apple stock around earnings periods.

The 10-day window was chosen, since we will be predicting if an extreme event occurs tomorrow based on the past 10 days of data.

```
[6]: from data_preprocessing import StockDataPreprocessor
# Instantiate the preprocessor and download the data
preprocessor = StockDataPreprocessor(ticker="AAPL", start_date="2015-01-01",
    end_date="2024-01-31")
stock_prices = preprocessor.download_and_prepare_stock_data()
# Add 10-day rolling volatility, volume relative to 10-day moving average and
    VIX index
stock_prices = preprocessor.add_features()
stock_prices.head()
```

```
[*****100%*****] 1 of 1 completed
/Users/nikolaosmourdoukoutas/ai2c_assignment/src/data_preprocessing.py:307:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self.data["relative_volume"] = self.data["Volume"] / volume_ma
[*****100%*****] 1 of 1 completed
```

```
[6]: Price          Adj Close      Close      High      Low  \
      Ticker          AAPL      AAPL      AAPL      AAPL
      Date
2015-02-27 00:00:00+00:00  28.720215  32.115002  32.642502  32.060001
2015-03-02 00:00:00+00:00  28.861069  32.272499  32.570000  32.075001
2015-03-03 00:00:00+00:00  28.921429  32.340000  32.380001  32.022499
2015-03-04 00:00:00+00:00  28.738096  32.134998  32.389999  32.080002
2015-03-05 00:00:00+00:00  28.261889  31.602501  32.187500  31.440001
```

Price Ticker Date	Open AAPL	Volume AAPL	Daily_Returns	Extreme_Event \
2015-02-27 00:00:00+00:00	32.500000	248059200	-1.502825	0.0
2015-03-02 00:00:00+00:00	32.312500	192386800	0.490435	0.0
2015-03-03 00:00:00+00:00	32.240002	151265200	0.209140	0.0
2015-03-04 00:00:00+00:00	32.275002	126665200	-0.633898	0.0
2015-03-05 00:00:00+00:00	32.145000	226068400	-1.657061	0.0

Price Ticker Date	rolling_volatility	relative_volume	VIX \
2015-02-27 00:00:00+00:00	0.014758	1.005357	13.34
2015-03-02 00:00:00+00:00	0.014758	0.787608	13.04
2015-03-03 00:00:00+00:00	0.014685	0.646066	13.86
2015-03-04 00:00:00+00:00	0.014717	0.553504	14.23
2015-03-05 00:00:00+00:00	0.015625	0.955874	14.04

Price Ticker Date	bollinger_band_width	ATR
2015-02-27 00:00:00+00:00	0.029195	0.59025
2015-03-02 00:00:00+00:00	0.026359	0.59900
2015-03-03 00:00:00+00:00	0.024543	0.58575
2015-03-04 00:00:00+00:00	0.024770	0.58350
2015-03-05 00:00:00+00:00	0.029292	0.64075

We will inspect the correlations between the features and the label. While this measure is not perfect, since it only captures linear relationships, it is still useful for indicating if the additional features are useful for the task at hand.

```
[7]: import numpy as np

correlations = stock_prices.corr()['Extreme_Event'].sort_values(ascending=False)
print("Correlations with Extreme_Event:")
print(correlations)
```

```
Correlations with Extreme_Event:
Price          Ticker
Extreme_Event      1.000000
VIX                0.310769
rolling_volatility 0.256255
bollinger_band_width 0.215680
ATR               0.213547
Volume            AAPL 0.192104
relative_volume    0.112583
High              AAPL 0.064290
```

```

Open          AAPL      0.062204
Close         AAPL      0.061895
Adj Close     AAPL      0.061462
Low           AAPL      0.059711
Daily_Returns      -0.055992
Name: Extreme_Event, dtype: float64

```

We see that the newly added features have a the highest correlation with the label, which is promising. We continue by refitting the random forest model with the new features.

The hyperparameter tuning and fitting of the model will be done in the same way as before. We will use the code defined in the module `src/random_forest.py`.

```

[8]: from random_forest import RandomForestOptimizer

# Split data into train, validation and test sets
X_train, y_train, X_val, y_val, X_test, y_test = StockDataPreprocessor.
    ↪split_data(
        stock_prices,
        ["Open", "High", "Low", "Close", "Volume", "rolling_volatility",
        ↪"relative_volume", "VIX", "bollinger_band_width", "ATR", "Daily_Returns"],
        "Extreme_Event",
        train_ratio=0.7,
        val_ratio=0.85,
    )

X_train, y_train = StockDataPreprocessor.time_series_to_supervised(
    X_train, y_train, lookahead=10
)
X_val, y_val = StockDataPreprocessor.time_series_to_supervised(X_val, y_val,
    ↪lookback=10)
X_test, y_test = StockDataPreprocessor.time_series_to_supervised(X_test,
    ↪y_test, lookahead=10)

# Initialize and run optimization
rf_optimizer = RandomForestOptimizer()
rf_optimizer.optimize(X_train, y_train, X_val, y_val)

# Train final model with best parameters
rf_model = rf_optimizer.train_best_model(X_train, y_train)

```

```

100%|      | 30/30 [01:02<00:00, 2.08s/trial, best loss:
-0.6553398058252428]

```

```

Best trial:
  Value (F2): 0.6553

```

```

Best parameters:
  max_depth: 79.0

```

```
max_features: 0.991728418919807
max_leaf_nodes: 582.0
min_impurity_decrease: 0.040044949236263346
min_samples_leaf: 3.0
min_samples_split: 9.0
n_estimators: 787.0
```

Let's evaluate the performance of the model on the test set.

```
[17]: from model_evaluation import ModelEvaluator


rf_evaluator = ModelEvaluator(
    model=rf_model,
    model_type="RF",
)

metrics = rf_evaluator.evaluate(X_test, y_test)
for metric, value in metrics.items():
    print(f"{metric:20s}: {value:.4f}")
print("-" * 40)

rf_evaluator.plot_confusion_matrix(X_test, y_test)
```

Seed set to 42

```
F2 Score           : 0.4885
F1 Score           : 0.3285
Precision          : 0.2125
Recall             : 0.7234
Accuracy           : 0.5749
AUC                : 0.6367
Balanced Accuracy  : 0.6367
```



The model seems to be doing significantly better than the random forest and the TCNN without the additional features, since it accurately predicts 34 of the 48 extreme events and 154 from the 283 non-extreme events.

The former achieves a test Area under the Receiver Operating Characteristic Curve of 0.63 approximately, while the latter yielded 0.5 and 0.52 respectively.

Let's plot the actual and predicted extreme events over time.

```
[18]: import matplotlib.pyplot as plt
import pandas as pd

y_pred = rf_model.predict(X_test.values)
y_pred = pd.Series(y_pred, index=y_test.index)

plt.figure(figsize=(15, 6))

# Plot actual events as stems
plt.stem(y_test.index[y_test == 1], y_test[y_test == 1],
        label='Actual Events', linefmt='b-', markerfmt='bo', basefmt=' ')

# Plot predicted events as red markers
plt.plot(y_test.index[y_pred == 1], y_pred[y_pred == 1],
        'r^', label='Predicted Events', markersize=8)
```

```

plt.title('Extreme Events: Actual vs Predicted (Random Forest + Additional_
↳Features)', pad=20)
plt.xlabel('Date')
plt.ylabel('Event Occurrence')

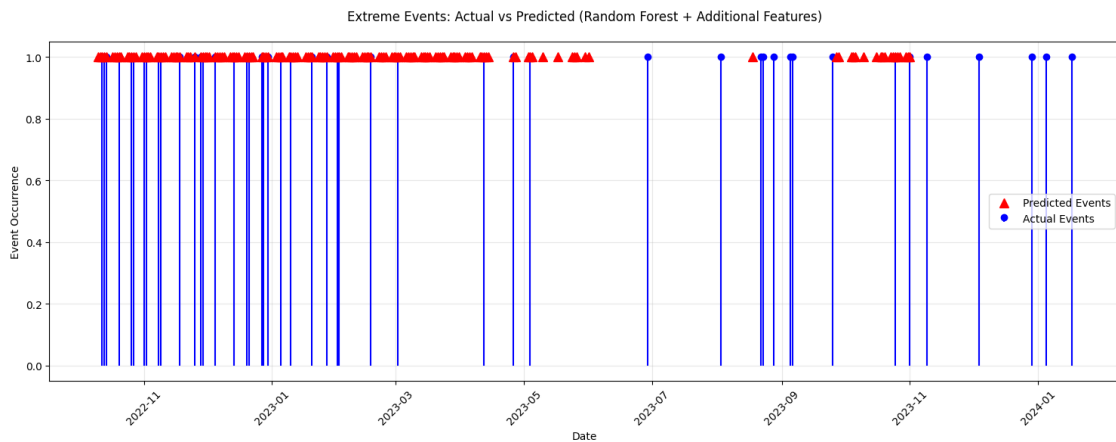
# Customize the plot
plt.grid(True, alpha=0.3)
plt.legend()

# Rotate x-axis labels for better readability
plt.xticks(rotation=45)

# Adjust layout to prevent label cutoff
plt.tight_layout()

# Save the plot
plt.savefig('../data/figures/rf_predictions_extra_features.png', dpi=300,
↳bbox_inches='tight')

```



The model tends to overpredict extreme events, particularly in late 2022 to early 2023, exhibiting high recall but lower precision in its predictions. This aligns with our rationale of prioritizing the capture of extreme events, even at the cost of false alarms. We observe that even though the frequency of extreme events drops after May 2023, the model predicts quite a few extreme events, especially between the September and November of 2023. A next step would be to explore strategies to reduce false positives while maintaining strong recall.

As a last step of analysis for the specific model, we will check the feature importance scores of the random forest model, which are indicative of the relevance of the features for the task at hand. These are computed as the mean decrease in impurity (MDI) provided by each feature when the tree is built. Specifically, the importance score is calculated by averaging the entropy reduction achieved by each feature at split points throughout all trees.

Let's check the top 10 features that have the highest importance scores.

```
[19]: importances = rf_model.feature_importances_
sorted_indices = np.argsort(importances)[::-1]
for i in sorted_indices[:10]:
    print(f"{X_train.columns[i]:20s}: {importances[i]:.4f}")
```

```
VIX(t-1)           : 0.8234
VIX(t-3)           : 0.0648
ATR(t-1)           : 0.0457
VIX(t-2)           : 0.0292
ATR(t-10)          : 0.0097
ATR(t-2)           : 0.0089
ATR(t-8)           : 0.0076
ATR(t-3)           : 0.0025
Volume(t-1)        : 0.0017
VIX(t-8)           : 0.0013
```

We see that the most important feature by far is the VIX index of the previous day, which is not surprising, since the VIX index is a measure of the market-wide volatility and extreme events are more likely during volatile periods. Overall VIX and ATR are the most important features, which is again expected, since they are measures of volatility.

These results could be useful for further feature engineering.

4.2 Long Short-Term Memory Model (LSTM)

We will now use a sequential neural network model. These architectures are well-suited for time series prediction tasks, since they are able to capture temporal dependencies and patterns in the data, which are inherent to time series.

Specifically, we will fit an LSTM model, which can handle long-term dependencies in the data, which are quite useful if trends across a large range of days influence the occurrence of an extreme event. On the contrary, the convolutional neural networks are limited to capturing local temporal patterns in the data, due to the fixed window size of the convolutional operations.

We will fit an unidirectional LSTM, since we are interested in predicting the occurrence of an extreme event tomorrow, given the past 10 days and bidirectionality would violate this temporal relationship and also result into data leakage. We don't expect a good performance from the LSTM, since our sample size is not very large and the model has a significant number of parameters.

```
[40]: # Instantiate the preprocessor and download the data
from data_preprocessing import StockDataPreprocessor
preprocessor = StockDataPreprocessor(ticker="AAPL", start_date="2015-01-01",
    end_date="2024-01-31")
stock_prices = preprocessor.download_and_prepare_stock_data()

# Add the extra features
stock_prices = preprocessor.add_features()
# Standardize the features
stock_prices = preprocessor.standardize_data(
    stock_prices, ["Open",
```

```

        "High",
        "Low",
        "Close",
        "Volume",
        "rolling_volatility",
        "relative_volume",
        "VIX",
        "bollinger_band_width"]
    )

# Split data into train, validation and test sets
X_train_lstm, y_train_lstm, X_val_lstm, y_val_lstm, X_test_lstm, y_test_lstm = 
    ↪ StockDataPreprocessor.split_data(
        stock_prices,
        ["Open",
         "High",
         "Low",
         "Close",
         "Volume",
         "rolling_volatility",
         "relative_volume",
         "VIX",
         "bollinger_band_width",
         "ATR",
         "Daily_Returns"],
        "Extreme_Event",
        train_ratio=0.7,
        val_ratio=0.85,
    )
# Convert the features into sequences
X_train_lstm, y_train_lstm = preprocessor.create_sequences(X_train_lstm, 
    ↪ y_train_lstm, lookback=10)
X_val_lstm, y_val_lstm = preprocessor.create_sequences(X_val_lstm, y_val_lstm, 
    ↪ lookback=10)
X_test_lstm, y_test_lstm = preprocessor.create_sequences(X_test_lstm, 
    ↪ y_test_lstm, lookback=10)

print(f"Training features shape: {X_train_lstm.shape}, labels shape: 
    ↪ {y_train_lstm.shape}")
print(f"Validation features shape: {X_val_lstm.shape}, labels shape: 
    ↪ {y_val_lstm.shape}")
print(f"Test features shape: {X_test_lstm.shape}, labels shape: {y_test_lstm.
    ↪ shape}")

```

```

[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed

```

Training features shape: (1561, 11, 10), labels shape: (1561,)
Validation features shape: (327, 11, 10), labels shape: (327,)
Test features shape: (327, 11, 10), labels shape: (327,)

The LSTM layer in Pytorch expects the input data to have the shape [samples, timesteps, features], so we need to permute the dimensions of the training, validation and test sets.

```
[41]: X_train_lstm = np.transpose(X_train_lstm, (0, 2, 1)) # From (samples, timesteps, features) to (samples, features, timesteps)
X_val_lstm = np.transpose(X_val_lstm, (0, 2, 1))
X_test_lstm = np.transpose(X_test_lstm, (0, 2, 1))

print(f"Training features shape: {X_train_lstm.shape}, labels shape: {y_train.shape}")
print(f"Validation features shape: {X_val_lstm.shape}, labels shape: {y_val.shape}")
print(f"Test features shape: {X_test_lstm.shape}, labels shape: {y_test.shape}")
```

Training features shape: (1561, 110), labels shape: (1561,)
Validation features shape: (327, 110), labels shape: (327,)
Test features shape: (327, 110), labels shape: (327,)

For the hyperparameter tuning we used Ray Tune package and specifically again the TPE BO algorithm. All this code is implemented in the `src/improvement.py` file. We load the best performing LSTM model and evaluate it on the test set.

```
[42]: from improvement import LSTM


best_lstm_config = {
    "batch_size": 16,
    "dropout_prob": 0.47163297383402925,
    "hidden_dim": 128,
    "learning_rate": 0.0019016102178314065,
    "num_layers": 3,
}

lstm_model = LSTM.load_from_checkpoint(
    "../data/models/best_lstm.ckpt",
    n_features=X_train_lstm.shape[2],
    lookback=X_train_lstm.shape[1],
    hidden_dim=best_lstm_config["hidden_dim"],
    num_layers=best_lstm_config["num_layers"],
    dropout_prob=best_lstm_config["dropout_prob"],
    learning_rate=best_lstm_config["learning_rate"],
)

lstm_evaluator = ModelEvaluator(model=lstm_model, model_type="LSTM")
lstm_metrics = lstm_evaluator.evaluate(X_test_lstm, y_test_lstm)
for metric, value in lstm_metrics.items():
    print(f"{metric:20s}: {value:.4f}")
print("-" * 40)
```

```
lstm_evaluator.plot_confusion_matrix(X_test_lstm, y_test_lstm)
```

F2 Score : 0.2778
F1 Score : 0.2796
Precision : 0.2826
Recall : 0.2766
Accuracy : 0.7951
AUC : 0.5794
Balanced Accuracy : 0.5794



We see that the LSTM model is definitely worse than the random forest model with the additional features, with an almost 13% drop in the balanced accuracy score and AUC. It is better than the TCNN, nevertheless, (57% AUC VS 52%) and thus we conclude that sequential architecture of the model was indeed beneficial.

Let's plot the actual and predicted extreme events over time.

```
[43]: y_pred = lstm_evaluator.predict(X_test_lstm)
      y_pred = pd.Series(y_pred, index=y_test.index)

      plt.figure(figsize=(15, 6))
```

```

# Plot actual events as stems
plt.stem(y_test.index[y_test == 1], y_test[y_test == 1],
        label='Actual Events', linefmt='b-', markerfmt='bo', basefmt=' ')

# Plot predicted events as red markers
plt.plot(y_test.index[y_pred == 1], y_pred[y_pred == 1],
        'r^', label='Predicted Events', markersize=8)

plt.title('Extreme Events: Actual vs Predicted (LSTM)', pad=20)
plt.xlabel('Date')
plt.ylabel('Event Occurrence')

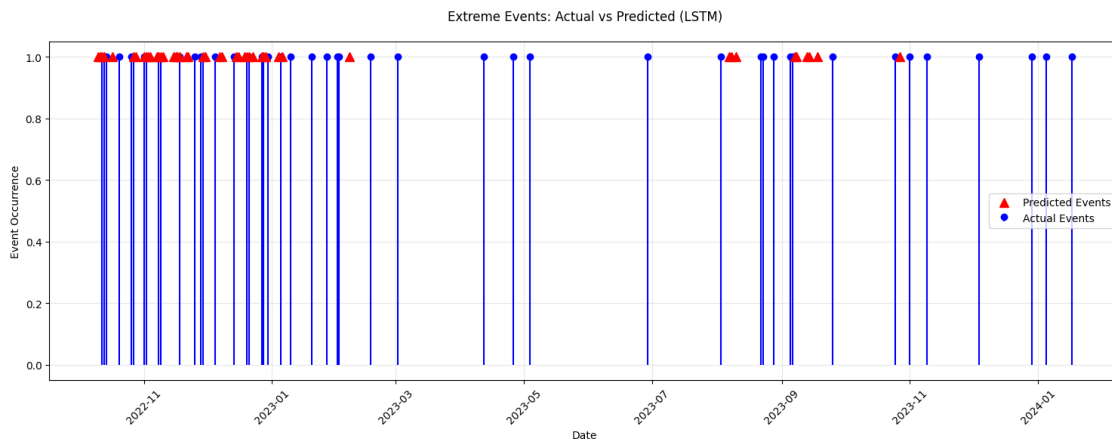
# Customize the plot
plt.grid(True, alpha=0.3)
plt.legend()

# Rotate x-axis labels for better readability
plt.xticks(rotation=45)

# Adjust layout to prevent label cutoff
plt.tight_layout()

# Save the plot
plt.savefig('../data/figures/lstm_predictions_extra_features.png', dpi=300,
        bbox_inches='tight')

```



The LSTM model shows a more conservative prediction behavior, correctly identifying only 13 extreme events while missing the vast majority of the rest. While the model achieves high precision with very few false positives, its low recall makes it unsuitable for our use case where we have assumed that capturing extreme events is prioritized over avoiding false alarms.

Also, after the hiatus of the extreme events, taking place after March 2023, it seems to be unable

to predict accurately the occurrence of extreme events, by contrast with the TCNN.

Note: A multilayer perceptron (MLP) model was also fitted, but due to its poor performance, it is not included in the report.

5 Conclusion

5.1 Model comparison

Let's compare the three best models: TCNN without extra features, random forest and LSTM with extra features. With respect to the classification metrics, the random forest outperforms the other two models, with the TCNN being the second best. The LSTM model is the worst performing model, with the lowest balanced accuracy and AUC scores.

But considering the task we need to consider the trade-off between precision and recall, as well as how their predictive performance evolves over time. Let's plot the cumulative accuracy of the three models over time.

```
[44]: from PIL import Image

fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(15, 20)) # Adjusted figsize
    ↪ for better vertical layout


# Load and display each image
img1 = Image.open('../data/figures/tcnn_predictions.png')
img2 = Image.open('../data/figures/rf_predictions_extra_features.png')
img3 = Image.open('../data/figures/lstm_predictions_extra_features.png')

ax1.imshow(img1)
ax1.axis('off')

ax2.imshow(img2)
ax2.axis('off')

ax3.imshow(img3)
ax3.axis('off')

plt.tight_layout()
plt.show()
```

The two best models are the random forest with the additional features and the TCNN. While TCNN is outperformed by the LSTM in terms of recall and the rest of the metrics, it seems that it is more accurate in predicting the extreme events towards the end of the time horizon.

The choice between the two models would depend on the risk appetite of the trader. If the trader

is risk averse, the TCNN would be the better choice, since it has a higher precision and a lower recall. If the trader is risk seeking, the random forest would be the better choice, since it has a higher recall and a lower precision.

5.2 Potential next steps for further improvements

One direction that could improve the performance of all the models is to consider additional features or construct new ones that are more informative for the task at hand. People with knowledge of the stock market could provide valuable insights on this topic.

Another direction would be to treat the class imbalance problem, by producing synthetic samples of extreme events. One methods that could be used is [T-SMOTE](#), a temporal version of SMOTE, designed specifically for handling imbalanced time series classification problems.

Finally, there are multiple works based on Generative Adversarial Networks (GANs) that are used for generating synthetic data of times series data, that could be used for increasing the count of extreme events in the training set. [TimeGAN](#) is one such method that could be used for this task.